

# Partial-order planning and hierarchical task networks

### Curriculum for 4:


- Section 11.3 *Partial-Order Planning* of **2nd edition** of Russell & Norvig (available on DTU Insider)
- Section 10.5 *Analysis of Planning Approaches* of Russell & Norvig 3ed
- Section 11.2 *Hierarchical Planning* of Russell & Norvig 3ed
- Jeff Orkin: *Three States and a Plan: The A.I. of F.E.A.R.* (available on DTU Inside).

Today's subjects:

- Partial-order planning.
- Hierarchical task networks.

## Serialisable sugoals

Planning problem with **serialisable subgoals**: There exists an order of the subgoals, such that they can be achieved in that order without destroying any of the previously achieved subgoals.




Trying the **Crunch level** from the warmup assignment:

Solver	Time	Plan length	Nodes generated
<i>BFS</i>	297s	98	$\approx 9 \cdot 10^6$
<i>BFS</i> utilising ser. subgoals	0.25s	106	27.831
<i>A*</i>	0.86s	103	24.044
<i>A*</i> utilising ser. subgoals	0.12s	106	411
greedy	0.25s	164	2.792
greedy utilising ser. subgoals	0.10s	104	214


## Linear sequences versus partial orders

Progression and regression planning explore **linear sequences of actions**.



This is often not advantageous. E.g. considering in which order to move the boxes will just make the planning process more time-consuming.

We'd like some degree of **partiality** in the ordering of actions, e.g. the order in which to move the boxes or the order in which to buy milk and bananas.



# Partial-order planning

A generalised treatment of partiality is found in **partial-order planning (POP)**.

**Partial-order planning (POP)**: A planning technique producing partially ordered plans, ignoring the ordering of independent actions.

Partial-order planning is in particular effective for **partly decomposable** problems like e.g. the gripper problem (or planning a party).

## Example: Socks-and-shoes

Putting on your socks and shoes...

### Action schemas:

ACTION: *LeftSock*

PRECONDITION:

EFFECT: *LeftSockOn*

ACTION: *LeftShoe*

PRECONDITION: *LeftSockOn*

EFFECT: *LeftShoeOn*

ACTION: *RightSock*

PRECONDITION:

EFFECT: *RightSockOn*

ACTION: *RightShoe*

PRECONDITION: *RightSockOn*

EFFECT: *RightShoeOn*

**Initial state:** empty.

**Goal:** *LeftShoeOn*  $\wedge$  *RightShoeOn*.

Note the independence of the actions on the left from the actions on the right. Partial-order planning takes advantage of this.

## Partial orders

In partial-order planning a **partial order** (order induced by directed, acyclic graph)  $\prec$  on the required actions is built up in a stepwise fashion.

**Example.** In the case of the socks-and-shoes problem, the partial order would become:

$$LeftSock \prec LeftShoe, RightSock \prec RightShoe.$$

Expresses: put on left sock **before** left shoe, and put on right sock **before** right shoe.


**Linearisation** of a partial order: Any total order that extends it.

Polynomial-time algorithms for linearisation: topological sort.

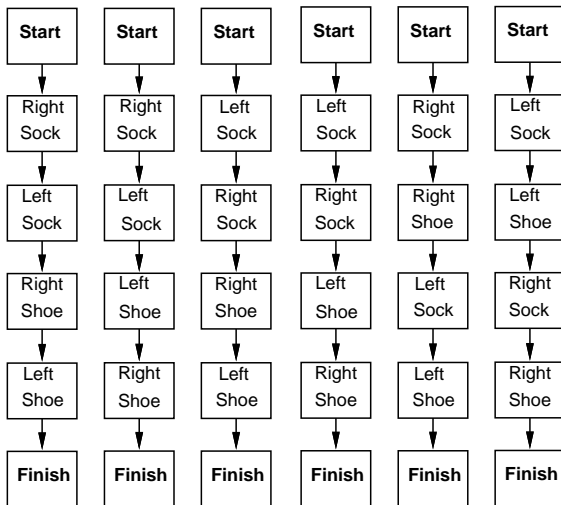
**Solution:** A partial order  $\prec$  is called a **solution** to a planning problem if any **linearisation** of it constitutes a **plan** (action sequence leading from start to goal).

# Partial-order planning for socks-and-shoes

Partial Order Plan:



Total Order Plans:



Compare with the simplified Gripper problem.

## Ordering constraints

**Ordering constraint:** An expression of the form  $A \prec B$ , where  $A$  and  $B$  are actions. Intended interpretation: Action  $A$  has to come before action  $B$ .

**Labelled ordering constraint:** An expression of the form  $A \overset{c}{\prec} B$ , where:

- $A$  and  $B$  are actions.
- $c$  is an **effect** of  $A$  (one of the literals in the effect of  $A$ ).
- $c$  is a **precondition** of  $B$  (one of the literals in the precondition of  $B$ ).

$A \overset{c}{\prec} B$  is read: “action  $A$  **achieves** (pre)condition  $c$  for performing action  $B$ ” or simply “ $A$  **achieves**  $c$  for  $B$ .” E.g.

$RightSock \overset{RightSockOn}{\prec} RightShoe.$

**Note:** Labelled ordering constraints are called **causal links** in Russell & Norvig (2ed), but I find it simpler to think of everything as ordering constraints.



# Partially ordered plans

**Partially ordered plan:** A directed graph where

- **Nodes** represent **actions**.
- **Edges** represent **(labelled) ordering constraints** between actions.

Ordering constraints:

*RightSock*  $\overset{RightSockOn}{\prec}$  *RightShoe*.

become edges:



**POP algorithm (partial-order planning algorithm):** Start with the empty partially ordered plan, then iteratively add more actions (nodes) and ordering constraints (edges).


# Open preconditions

Recall:  $A \stackrel{c}{\prec} B$  reads “ $A$  **achieves**  $c$  for  $B$ .”

Let  $B$  denote an action occurring in a partially ordered plan, and let  $c$  denote one of its preconditions (precondition literals).

**Open precondition:** The precondition  $c$  is called **open** if it is not achieved by any action in the plan, that is, the plan doesn't contain any ordering constraint of the form  $A \stackrel{c}{\prec} B$ .

**Example.** In the partially ordered plan below, *RightSockOn* is open, but not *LeftSockOn*.



# Conflicts

**Conflict:** An action  $D$  is said to **conflict** with an ordering constraint  $A \overset{c}{\prec} B$  if  $D$ 's effect is inconsistent with  $c$  ( $D$  destroys  $c$ ). This means that we cannot allow action  $D$  to be carried out between carrying out  $A$  and carrying out  $B$ .

**Example.** The action  $Go(N, H)$  conflicts with the ordering constraint  $Go(H, N) \overset{At(N)}{\prec} Buy(M, N)$ : we are not allowed to perform the action  $Go(N, H)$  between the actions  $Go(H, N)$  and  $Buy(M, N)$ .

**Resolving a conflict:** If  $D$  conflicts with an ordering constraint  $A \overset{c}{\prec} B$  we can **resolve** the conflict by adding either the constraint  $B \prec D$  or  $D \prec A$ .



# POP algorithm

**POP algorithm:** Start with the empty plan. Then iteratively do the following:

- Pick an **open precondition** and choose an action that achieves it.
- Resolve any conflicts introduced.

**Solution:** A partially ordered plan with no open preconditions and no unresolved conflicts will constitute a **solution** to the planning problem (because any linearisation of the ordering  $\prec$  of the partially ordered plan will constitute a plan).

Now for the details...

## POP algorithm

**Initialisation:** One starts with the **empty plan** only consisting of the actions *Start* and *Finish*.

*Start:* An action having the planning problem's start state description as effect.

*Finish:* An action having the planning problem's goal state description as precondition.

After initialisation the following step is repeated:

**Step:** Pick an open precondition  $c$  on some action  $B$  in the current plan and choose an action  $A$  that achieves  $c$  (either an existing action in the plan or a new one). Then:


1. Add ordering constraints  $A \overset{c}{\prec} B, Start \prec A, A \prec Finish$  (not necessary to draw edges for the two latter).
2. Resolve any conflicts introduced.

**Backtrack** if:

- An open precondition can't be achieved.
- Some conflict can't be resolved.

## POP example: simplified Gripper problem

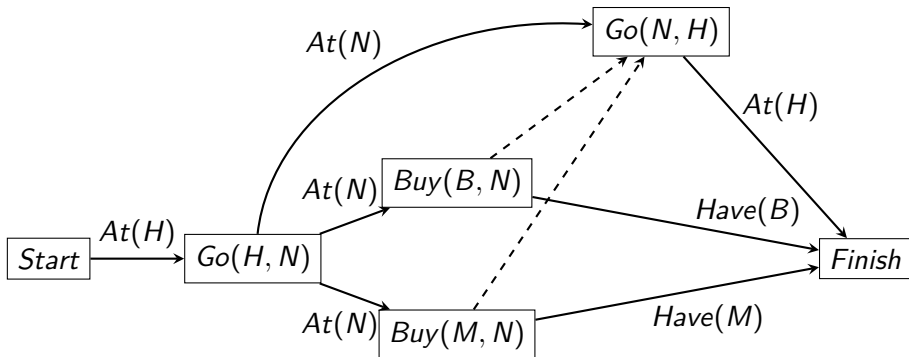
The following is the partially ordered plan resulting from running the POP algorithm on the simplified Gripper problem with  $n = 3$ :



Note that the graph constructed only has size  $O(n)$ , even less than the  $O(n^2)$  for greedy-best first GRAPH-SEARCH with an optimal heuristics.

## POP example: milk-and-bananas

Start is  $At(H)$ . Goal is  $At(H) \wedge Have(M) \wedge Have(B)$ . Partially ordered plan produced by POP algorithm (dashed edges are conflict resolutions):



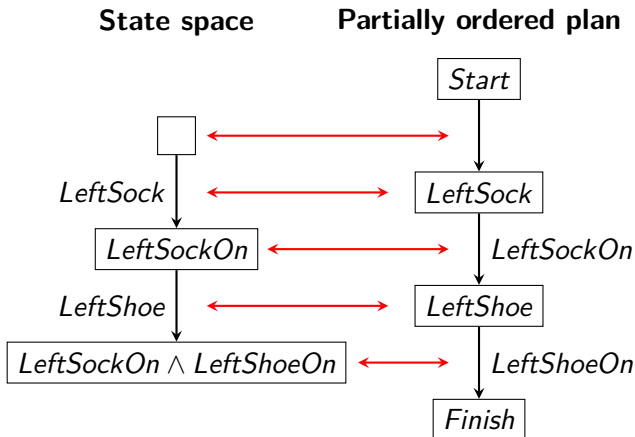
The linearisations are  $Go(H, N), Buy(B, N), Buy(M, N), Go(N, H)$  and  $Go(H, N), Buy(M, N), Buy(B, N), Go(N, H)$ .

Is partial-order planning closer to the way humans make plans? Probably. We tend to think in terms of “open preconditions” and postpone choices of ordering of independent actions.

# Duality

Note that partially ordered plans are **dual** to state space graphs! (Nodes in one is edges in the other.)

**Example.** Initial state is empty, goal is *LeftShoeOn*.





# POP Summary

## Properties of the **POP** algorithm:

- Works from goal to start state (as regression planning and the additive/max heuristics).
- Builds a graph (partial order plan) dual to a state space graph.
- Any linearisation of the partial order will constitute a solution to the planning problem.
- **Advantages:**
  - Particularly useful for partly decomposable problems. **Why? Because one doesn't have to decide on the ordering of independent actions.**
  - Good for **replanning** in dynamic worlds.
  - Extends naturally to **planning and scheduling problems**.
  - Generates human-friendly plans.
  - Important for applications (e.g. Mars rovers, Maersk container shipping).
- **Disadvantage:** Doesn't currently scale up as well to large problems as state-space planners with good heuristics (hard to come up with good heuristics for POP).

# Hierarchical planning

Partial-order planning explores **problem decomposition** to gain efficiency: Independent subgoals are worked on independently.

An alternative problem decomposition is **hierarchical decomposition**: Split **complex tasks** into a small number of **subtasks** that each in turn consists of a small number of subtasks, etc.

Humans have a similar approach to planning.

**Example.** Assume my goal is to have lunch.

- Top-level plan: *FindFood, EatFood*.
- *FindFood* can be split into *ChooseShop, BuyFoodAtShop*.
- *BuyFoodAtShop* can e.g. be split into *GoToShop, BuyFood, GoHome*.

Contrast this with basic planning at the atomic level of, say, individual body movements.

# HTN planning

Hierarchical planning can reduce the combinatorial explosion resulting from planning at the atomic level.

**Hierarchical Task Network planning (HTN planning):** A planning method exploring the ideas of hierarchical planning.

The **essential component** of HTN planning is **action decompositions**.

**Action decompositions (or refinements):** The reduction of a **high-level action (HLA)** to a sequence of lower-level actions.

**HTN planning algorithm:** HLA's are recursively refined until only **primitive actions** (non-refinable actions) remain.

**Implementation** of an HLA: The result of a recursive refinement into primitive actions. Usually not unique.

# HTN refinements

HTN **refinements** are represented by **refinement schemas**. A **refinement schema** consists of the following elements:

- **HLA name.** The name of the HLA that the refinement schema defines a decomposition for. E.g. *GetFood*, *Buy2Things*( $x, y$ ) or *Navigate*( $x, y$ ).
- **Precondition.** A conjunction of function-free literals. E.g.  $At(Shop) \wedge Sells(Shop, x) \wedge Sells(Shop, y)$ . Tells which conditions have to be satisfied in order for the refinement to be applicable.
- **Steps.** A sequence of actions (a totally ordered plan). These actions can either be primitive or high-level. E.g.  $[Buy(x, Shop), Buy(y, Shop)]$  or  $[Go(x, z, r), Navigate(z, y)]$ .

## Example.

**REFINEMENT:** *Buy2Things*( $x, y$ )

**PRECONDITION:**  $At(Shop) \wedge Sells(Shop, x) \wedge Sells(Shop, y)$

**STEPS:**  $[Buy(x, Shop), Buy(y, Shop)]$

# Refinement properties

- **Precondition.** Sometimes the **precondition** of a refinement schema is not explicitly stated, as it can be computed from the **steps** of the schema. **How?** All open preconditions when considering the steps as a partially ordered plan with no start state.
- **Multiple refinements.** There can be several refinements schemas for the same high-level action. E.g. a refinement of  $Navigate(x, y)$  with steps  $[Go(x, z, r), Navigate(z, y)]$  and another with just a single step  $[Go(x, y, r)]$ . Note the recursion!

**High-level plan:** A sequence of high-level and possibly primitive actions.

**Achieving a goal:** A high-level plan **achives a goal** if *at least one* of its implementations (recursive refinements) does.

# HTN planning example

## Example. Refinement schemas:

**REFINEMENT:**  $Navigate(x, y)$

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(z, r)$


**STEPS:**  $[Go(x, z, r), Navigate(z, y)]$

**REFINEMENT:**  $Navigate(x, y)$


**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r)$

**STEPS:**  $[Go(x, y, r)]$

## Initial state:



Assume Shakey is asked to carry out the HLA  $Navigate(L1, L3)$ ....



**Implementation** of the HLA *Navigate*( $L1, L3$ ):

1. Refine *Navigate*( $L1, L3$ ) into

$[Go(L1, D1, Room1), Navigate(D1, L3)]$ .

2. Then further into:

$[Go(L1, D1, Room1), Go(D1, D2, Room2), Navigate(D2, L3)]$ .

3. And finally into:

$[Go(L1, D1, Room1), Go(D1, D2, Room2), Go(D2, L3, Room3)]$ .

*Navigate*( $L1, L3$ ) **achieves the goal** *At*( $L3$ ).

## Representing high-level actions

To allow planning in terms of HLA's before refinement, HLA's are specified using PDDL/STRIPS **action schemas** exactly as primitive actions are.

**Example.** The  $Navigate(x, y)$  HLA can be represented by the following PDDL/STRIPS **action schema**:

**ACTION:**  $Navigate(x, y)$

**PRECONDITION:**  $At(x)$

**EFFECT:**  $At(y) \wedge \neg At(x)$

**Downward refinement property** of an HLA action schema: At least one implementation of the HLA achieves the stated **EFFECT** of the schema from the stated **PRECONDITION**.

**Example.** The HLA action schema above has the downward refinement property wrt. the refinements defined earlier.

**Note:** We are here presenting a simplified version of HTN where the effects of a high-level action are unique, not using **angelic semantics** as in Russell & Norvig 3ed.



## HTN planning example

**Example.** Assume now that the robot can only navigate through a room if the light is on in the room:

**ACTION:**  $Go(x, y, r)$

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r) \wedge Light(r)$

**EFFECT:**  $At(y) \wedge \neg At(x)$

This modification causes a change in the preconditions of the  $Navigate(x, y)$ -refinements:

**REFINEMENT:**  $Navigate(x, y)$

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(z, r) \wedge Light(r)$

**STEPS:**  $[Go(x, z, r), Navigate(z, y)]$

**REFINEMENT:**  $Navigate(x, y)$

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r) \wedge Light(r)$

**STEPS:**  $[Go(x, y, r)]$

## HTN planning example (cont'd)

Recall the action schema of  $Navigate(x, y)$ :

**ACTION:**  $Navigate(x, y)$

**PRECONDITION:**  $At(x)$

**EFFECT:**  $At(y) \wedge \neg At(x)$

With the refinement modifications carried out above,  $Navigate(x, y)$  unfortunately no longer satisfies the downward refinement property.

**Why?** If the light is off, no navigation is possible.

**Naive solution:** Add to the precondition of the action schema of  $Navigate(x, y)$  that the light should be on in all rooms. This is no good.

**Why?** Impractical and only works if the domain (number and name of rooms) is fixed.

**Can you think of a better solution?** Redefine refinements to allow turning on the light. See following slides.

## HTN planning example (cont'd)

**Better solution:** Replace  $Go(x, y, r)$  in the steps of the  $Navigate(x, y)$ -refinements by the following HLA  $Go'(x, y, r)$ :

**ACTION:**  $Go'(x, y, r)$

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r)$

**EFFECT:**  $At(y) \wedge \neg At(x) \wedge Light(r)$

**REFINEMENT:**  $Go'(x, y, r)$

**PRECONDITION:**

$At(x) \wedge In(x, r) \wedge In(y, r) \wedge Light(r)$

**STEPS:**  $[Go(x, y, r)]$

**REFINEMENT:**  $Go'(x, y, r)$

**PRECONDITION:**

$At(x) \wedge In(x, r) \wedge In(y, r)$

**STEPS:**  $[TurnOn(x, r), Go(x, y, r)]$

Here  $TurnOn(x, r)$  is a **primitive action** for turning on the light in the current room:

**ACTION:**  $TurnOn(x, r)$

**PRECONDITION:**  $At(x) \wedge In(x, r)$

**EFFECT:**  $Light(r)$

**Downward refinement property re-established:** Light doesn't have to be on for  $Navigate(x, y)$  to be executable.

## HTN planning example (cont'd)

We now have the following 3-level action hierarchy (primitive actions not shown):

**REFINEMENT:** *Navigate*( $x, y$ )

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(z, r)$

**STEPS:** [*Go'*( $x, z, r$ ), *Navigate*( $z, y$ )]

**REFINEMENT:** *Navigate*( $x, y$ )

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r)$

**STEPS:** [*Go'*( $x, y, r$ )]

**REFINEMENT:** *Go'*( $x, y, r$ )

**PRECONDITION:**

$At(x) \wedge In(x, r) \wedge In(y, r) \wedge Light(r)$

**STEPS:** [*Go*( $x, y, r$ )]

**REFINEMENT:** *Go'*( $x, y, r$ )

**PRECONDITION:**  $At(x) \wedge In(x, r) \wedge In(y, r)$

**STEPS:** [*TurnOn*( $x, r$ ), *Go*( $x, y, r$ )]

**ACTION:** *Navigate*( $x, y$ )

**PRECONDITION:**  $At(x)$

**EFFECT:**  $At(y) \wedge \neg At(x)$

**ACTION:** *Go'*( $x, y, r$ )

**PRECONDITION:**

$At(x) \wedge In(x, r) \wedge In(y, r)$

**EFFECT:**

$At(y) \wedge \neg At(x) \wedge Light(r)$

Note the increase in **planning efficiency**: when planning the route, no branching is produced from considering whether the light is on or not. This problem is dealt with *after* the overall route has been planned.

## HTN planning example (cont'd)

We can add further levels on top of our existing plan hierarchy:

**ACTION:**  $MoveBox(b, x, y)$

**PRECONDITION:**  $On(b, x)$

**EFFECT:**  $On(b, y) \wedge \neg On(b, x)$

**REFINEMENT:**  $MoveBox(b, x, y)$

**PRECONDITION:**  $At(l) \wedge On(b, x)$

**STEPS:**  $[Navigate(l, x), NavWithBox(b, x, y), Navigate(y, l)]$

**ACTION:**  $NavWithBox(b, x, y)$

**PRECONDITION:**  $At(x) \wedge On(b, x)$

**EFFECT:**  $At(y) \wedge \neg At(x) \wedge On(b, y) \wedge \neg On(b, x)$

**REFINEMENT:**  $NavWithBox(b, x, y)$

**PRECONDITION:**  $At(x) \wedge On(b, x) \wedge In(x, r) \wedge In(z, r)$

**STEPS:**  $[Push'(b, x, z, r), NavWithBox(b, z, y)]$


**REFINEMENT:**  $NavWithBox(b, x, y)$

**PRECONDITION:**  $At(x) \wedge On(b, x) \wedge In(x, r) \wedge In(y, r)$

**STEPS:**  $[Push'(b, x, y, r)]$

## HTN planning example (cont'd)

We now have an HLA  $MoveBox(b, x, y)$  for moving arbitrary boxes between arbitrary locations. Consider the following scenario:



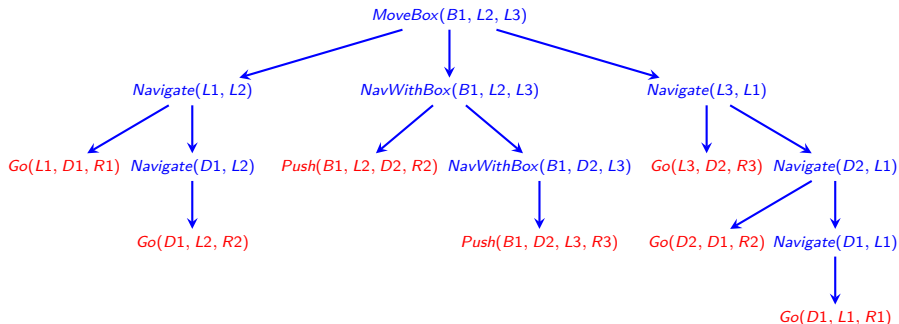

Assume our **goal** is  $On(B1, L3)$ .

The HLA  $MoveBox(B1, L2, L3)$  **achieves** this goal.

The following slide shows a possible implementation.

**Remark.** For simplicity, we ignore the light aspect (thus using  $Go$  instead of  $Go'$  and  $Push$  instead of  $Push'$ ).


# Implementation example



Note how the hierarchical approach saves on search/branching.

# HTN planning advantages

**Huge advantage** of HTN planning: Can be arbitrarily adjusted along the axis from **purely imperative** (no search) to **purely declarative** (only search):



Allows for combination of subtasks of different types:

- Trivial ones with unique (or few) implementations. (Solution is given).
- Subtasks corresponding to classical planning problems. (Solution completely unknown).



# HTN planning advantages

## Advantages of HTN planning:

- Allowing **hierarchical decomposition**, an approach widely used in “real life” (consider e.g. writing a large report or planning your vacation).
- “Help” the planner by telling how complex tasks can be decomposed. Can be the crucial difference between **intractability** and **tractability** for **large-scale applications**.

HTN planning has been more popular in **applications** and **industry** than any other planning method:

- Used for production-line scheduling, spacecraft planning and scheduling, equipment configuration, manufacturing process planning, evacuation planning, bridge computers, robotics.
- Used by companies such as Hitachi, Price Waterhouse, and Jaguar Cars.

## Final remark

Remember to read the following article (available on CampusNet):

- Jeff Orkin: “Three States and a Plan: The A.I. of F.E.A.R.”. Game Developers Conference (GDC), 2006.

It's an entertaining and fairly easy read, it presents a nice planning application, and provides a good understanding of some of the strengths and weaknesses of classical planning.