

02285 AI and MAS, SP18

Multiagent planning and agent architectures

Curriculum for week 5: Sections 11.3.3 *Online replanning* and 11.4 *Multiagent planning* of Russell & Norvig 3ed + Chapter 2 of Michael Wooldridge's "Reasoning about rational agents" (available on Campus-Net).

Today's subjects:

- Online planning.
- Multiagent planning.
- Agent architectures.
- The BDI agent architecture.

Programming project solution example 2016

http://www2.compute.dtu.dk/~tobo/Almuffins_on_MAsojourner.mov

Assumptions in classical planning

Assumptions in classical planning (PDDL):

1. Environments are **deterministic**.
2. Environments are **fully observable**.
3. Environments are **static** (no exogenous events).
4. Planning is **offline** (compute entire plan before starting to act).
5. Environments are **single-agent** (only one planning agent).

Today we will consider loosening assumptions 3, 4 and 5 (the most relevant to the programming project).

Offline planning vs. online replanning

Offline planning: Compute entire plan before executing first action (plan first, act later).

- **Ideal** when the domain is: **deterministic, static, fully observable**.
- **Problematic** in all other cases: one has to plan for **all contingencies** which easily leads to a combinatorial explosion.

Example. Box locations in hospital domain not known in advance.

A **better solution** is often to do **online (re)planning**: interleave planning and acting.

Even humans do seldom do pure offline planning, for the same reasons that AI agents shouldn't: planning for all eventualities can lead to unmanageable branching factors.

The necessity of replanning

Online replanning is in some cases our only option:

- The agent's model of the world (action schemas) might have **missing effects** or even **incorrect effects**. E.g. Shakey might leave a trace of oil behind itself when moving; or the effect of *TurnOn*(x, r) might fail to appear if the light bulb in room r is blown.
- More generally, action descriptions can be **incorrect** in any number of ways, and only a replanning (and/or learning) agent can be robust against such problems.
- It can be impossible for the agent to model all possible **exogenous events** and their effects, e.g. unexpected phone calls or earthquakes.

Execution monitoring

Execution monitoring: Monitoring the execution of a computed plan in the case of online planning. Used to figure out when **replanning** is required.

Three levels of **execution monitoring** (from simple to more advanced):

- **Action monitoring:** when an action is about to be performed, check that all preconditions of the action still hold (e.g. by **sensing**).
- **Plan monitoring:** when an action is about to be performed, check that the remaining plan will still succeed (nothing has destroyed the possibility of carrying out the remaining actions of the plan).
- **Goal monitoring:** when an action is about to be performed, check to see if there is a better set of goals that the agent might try to achieve. This requires some degree of **intentionality**, which will be discussed further in the following.

Online replanning

Challenge: How to avoid that **replanning** means planning all over from scratch?

E.g. **partial-order planning** techniques are quite good, as they allow significant reuse in case of replanning.

Example. (plan repair). Simple approach: whenever action monitoring calls for replanning, **repair** the plan by planning how to get back to a state from which the original plan can be continued.

Sometimes too naive: consider the case of the sphex wasp.



Multiagent planning

Multiagent planning problem: Planning problem with several planning agents.

Goals: Can be **individual**, **common** (e.g. players on the same team) or **conflicting** (e.g. players on opposing teams).

Basic multiagent concepts:

- **(Intelligent) agent:** Computer system/process capable of **independent (autonomous)** action in some environment—pursuing its own goals (**goal-directed behavior**).
- **Proactive agent:** Generating and seeking to achieve goals; not driven solely by external events; taking the **initiative**.
- **Multiagent system:** A system of **interacting (proactive) agents**. E.g. systems of hospital robots, softbots on the Internet, teams of soccer-playing robots or non-player characters in video games.
- Agents sharing the same environment (resources) will most often need to **cooperate**, **coordinate**, and **negotiate** with each other. Cf. human social interaction.

Single-agent multibody planning

Multibody planning problem: A single agent is planning on behalf of all agents. **Centralised** planning.

Multibody planning essentially reduces to single-agent planning (the other agents are reduced to “puppets”).

What could be the advantages/disadvantages of multibody planning (centralised) compared to “true” multiagent planning (decentralised)?

Multiagent architecture: The arrangement of the agents and their protocols of interaction and collaboration in a multiagent system.

Example: A centralised planning agent distributing goals (jobs) to the other agents that then compute their own individual plans for achieving these goals.

Example: The Martha project

The **Martha project** (1998) studies a planning system for dock-worker robots.

Multiagent architecture of Martha:


1. A **centralized planner** distributes transportation jobs to the individual dock-worker robots.
2. The robots then plan and schedule their own transportation jobs **individually**, not taking the other robots into account.
3. Each computed individual plan is **publicly broadcasted**.
4. In case of **resource conflicts**, the robot causing it will be asked to replan, see next slide.

Example cont'd

Resolving resource conflicts: Robot r broadcasts its computed plan. If robot s needs one of r 's scheduled resources, s will ask r to replan.

Problem: r might not be able to replan, e.g. in case of a **deadlock loop**.

In this case, the robot r will perform **centralised planning (multibody planning)** for the subset of robots involved in the conflict/deadlock loop.



Centralised vs. decentralised planning

Decentralised systems lead to a number of non-trivial challenges:

- **Communication.**
- **Coordination.**
- **Deadlock handling.**
- ...

So isn't a centralised system (multibody planning) the better choice?

No, the complexity of multibody planning **grows exponentially** with the number of agents planned for...

Example. Suppose a centralised planner plans on behalf of all dock-worker robots. It has to consider which **joint actions** $\langle a_1, \dots, a_n \rangle$ involving the n robots has to be performed. If each robot has b different possible actions to perform, the branching factor of the search becomes b^n . Note the number of agents in the exponent!

When the agents' tasks are **loosely coupled** as in the Martha project, this is unnecessarily expensive.

Example cont'd

In the Martha project, the distribution of planning tasks reduces the complexity to **almost linear** in the number of robots (up to at least 50 robots in total).

Distributing planning tasks between agents is an example of **problem decomposition**. Relates to **partial-order planning** and **hierarchical planning**.

General approach underlying Martha project:

- First “pretend” the individual agents’ problems are completely decoupled (a **relaxed problem**).
- Solve the individual problems separately.
- Fix up the interactions (**conflicts**) by communication and coordination.

Joint plans in multiagent planning

Consider the **milk-and-bananas problem** involving two agents, $A(nn)$ and $B(ob)$. The common goal of the two agents is to have milk and bananas.

A **joint plan** that would work (assuming **perfect synchronization**):

$A : Go(H, N), Buy(M), Buy(B), Go(N, H)$

$B : WatchTV, WatchTV, WatchTV, WatchTV$

Ann, however, might have something more along the following lines in mind:

$A : WatchTV, WatchTV, WatchTV, WatchTV$

$B : Go(H, N), Buy(M), Buy(B), Go(N, H)$

Deciding which joint plan to choose requires either **communication** (“you do it, Bob!”) or some kind of agreed **convention** or **social law** (Bob always does the shopping on Tuesdays).

Example. Collision avoidance when driving vs. when walking.

Task sharing in multiagent planning

The **optimal plan** for Ann and Bob in terms of number of joint actions is:

$A : Go(H, N), Buy(M), Go(N, H)$

$B : Go(H, N), Buy(B), Go(N, H)$

It might be achieved through **negotiation** (“I’ll go if you come along”).

Task sharing: How to share tasks between agents through **communication** and **coordination**.

Common methods for **task sharing**:

- **Negotiation.**
- **Bidding.**
- **Voting.**
- **Conventions and social laws.**

Multibody planning

All single-agent planning algorithms presented in this course can with only minor modifications be used for centralised **multibody planning**.

One should however be aware of the potential **exponential blowup** of complexity in the number of agents.

Partial-order planning can work well in multibody planners for loosely coupled multiagent problems. **Why?**

Agent architectures

In complex dynamic multiagent environments, agents have to be more than just planners.

We must consider more complex **architectures** for agents. In these, agents are build from many independent **computational modules**, of which the planning engine is just one.

Definition of **agent architecture** from Patti Maes (1991):

A particular methodology for building agents. It specifies how [...] the agent can be decomposed into the construction of a set of component modules and how these modules should be made to interact. [...] An agent architecture encompasses techniques and algorithms that support this methodology.

One such agent architecture is the **BDI architecture**, that we will describe in the following.

Deliberation and means-end reasoning

In pure planning, an agent is given a **predefined goal** to be achieved.

Planners are thus suitable for finding out **how** to do something, but not for finding out **what** to do.

- **Deliberation:** The process of finding out **what** state of affairs (which goal) we want to achieve.
- **Means-end reasoning:** The process of deciding **how** to achieve this state of affairs

A planner does means-end reasoning but not deliberation.

A **proactive agent** will both need a deliberation module and a means-end reasoning module (planner).

Intentions and the intentional stance

- The output of a **deliberation process** is **intentions**: the state of affairs that the agent has **committed** itself to achieving. E.g. buying milk and bananas.
- The output of a **means-end reasoning process** is a **plan**: how to achieve the intentions. E.g. a plan for buying milk and bananas.

The intentional stance: Describing agents in terms of their beliefs, intentions, desires, hopes, etc.

The intentional stance is a very **powerful abstraction**: It provides us with a convenient and familiar way of describing, explaining, and predicting the behaviour of complex systems.

It also helps in **designing** such AI systems (in analogy to how the **object** abstraction helps us designing large computer programs).

Example. Taking the intentional stance towards a thermostat will perhaps not buy us much, but e.g. in the dock-worker environment it will.

Intentions

Since **intentions** are goals that an agent commits to, if it has the intention ϕ then usually the following should be the case:

- The agent should **devote resources** to obtaining ϕ .
- The agent should not adopt an intention in **conflict** with ϕ (e.g. eat the banana and drink the milk on the way back from Netto).
- **Persistence**: If the agent's first attempt on achieving ϕ fails, it should **try again** (possibly using a different plan).
- The agent should **believe** that ϕ is possible to achieve, and that under “normal circumstances” it will in fact eventually achieve ϕ .

Intentions vs. desires

Note that intentions are stronger than desires: **desires** are everything we would like to achieve in an “ideal world” (be young, rich, famous, beautiful, happy, etc.), and some of these might be unattainable or mutually exclusive.

In contrast, **intentions** are desires that we have already chosen to commit to (e.g. complete ones studies in the hope that this is a step on the way to become rich).

The BDI architecture: version 1

The **BDI architecture** is an agent architecture based on the 3 attitudes **belief**, **desire**, and **intention**. An agent designed by the BDI architecture will continuously be engaged in an **agent control loop** where the agent alternately does deliberation and means-end reasoning.

First version of the **agent control loop**:

Agent Control Loop Version 1

1. while true do
2. **perceive** the world;
3. update **beliefs** about the world;
4. **deliberate** about what **intention** to achieve next;
5. use **means-ends reasoning** to compute
 a **plan** for the intention;
6. **execute** the plan
7. end while

Steps 5 and 6 corresponds to what we have already been covering on automated planning.

BDI formal notation

We will now make the agent control loop a bit more formal. To be able to do this, we must first introduce some new notation:

- **B : current beliefs.** $B \in Bel$: the set of **all possible beliefs**.
- **D : current desires.** $D \in Des$: the set of **all possible desires**.
- **I : current intentions.** $I \in Int$: the set of **all possible intentions**.
- **ρ : a percepts.** $\rho \in Per$: the set of **all possible percepts**.
- **π : plan.** $\pi \in Plan$: the set of **all possible plans**. Plans are sequences of actions.
- **$hd(\pi)$: first action** of the plan π .
- **$tail(\pi)$: remaining sequence of actions** in π .
- **$execute(\pi)$: the procedure** executing π .

BDI formal notation (cont'd)

- **Belief revision function:**

$$brf : \mathbb{P}(Bel) \times Per \rightarrow \mathbb{P}(Bel)$$

$brf(B, \rho)$ is your updated beliefs based on current beliefs B and percept ρ .

- **Deliberation function:**

$$deliberate : \mathbb{P}(Bel) \rightarrow \mathbb{P}(Int)$$

$deliberate(B)$ is the set of intentions chosen based on current beliefs B .

- **Planning function:**

$$plan : \mathbb{P}(Bel) \times \mathbb{P}(Int) \rightarrow Plan$$

$plan(B, I)$ is the plan chosen based on current beliefs B and current intentions I .

The BDI architecture: version 2

Using the notations introduced above, we can now express the agent control loop in a slightly more formal way:

Agent Control Loop Version 2

1. $B := B_0$ /* initial beliefs */
2. while true do
3. get next percept ρ ;
4. $B := brf(B, \rho)$;
5. $I := deliberate(B)$;
6. $\pi := plan(B, I)$;
7. execute(π)
8. end while

Deliberation

The deliberation step (step 5 in the algorithm above) often splits into two parts:

1. Generate the **desires**, that is, everything the agent might want to achieve.
2. Generate the **intentions** by choosing between competing desires and commit to some of them.

First part is called **option generation**:

$$options : \mathbb{P}(Bel) \times \mathbb{P}(Int) \rightarrow \mathbb{P}(Des)$$

Second part is called **filtering**:

$$filter : \mathbb{P}(Bel) \times \mathbb{P}(Des) \times \mathbb{P}(Int) \rightarrow \mathbb{P}(Int)$$

The BDI architecture: version 3

Splitting the **deliberation step** into an **option generation step** followed by a **filtering step**, we now get the following agent loop:

Agent Control Loop Version 3

1. $B := B_0$ /* initial beliefs */
2. $I := I_0$ /* initial intentions */
3. while true do
4. get next percept ρ ;
5. $B := brf(B, \rho)$;
6. $D := options(B, I)$;
7. $I := filter(B, D, I)$;
8. $\pi := plan(B, I)$;
9. execute(π)
10. end while

Blind commitment vs. plan monitoring

Weakness of current agent control loop: The agent is **blindly committed** to both its **ends** (intentions) and its **means** (plans).

In other words, the agent is doing **offline deliberation** and **offline planning**.

This is insufficient if the domain is **partially observable**, **dynamic** (e.g. has other agents) or **nondeterministic**.

First step towards **online (re)planning** is to introduce **plan monitoring**: **Replan** if current plan is no longer **sound** given the agent's intentions and beliefs.

We use the notation $sound(\pi, I, B)$ to mean that π is a **sound plan** for achieving I given beliefs B .

Version 4 of the agent control loop integrates this type of **plan monitoring**...

The BDI architecture: version 4

Agent Control Loop Version 4

```
1.  $B := B_0; I := I_0;$ 
2. while true do
3.   get next percept  $\rho$ ;
4.    $B := brf(B, \rho); D := options(B, I); I := filter(B, D, I);$ 
5.    $\pi := plan(B, I);$ 
6.   while  $\neg empty(\pi)$  do
7.      $execute(hd(\pi));$ 
8.      $\pi := tail(\pi);$ 
9.     get next percept  $\rho$ ;
10.     $B := brf(B, \rho);$ 
11.    if not sound( $\pi, I, B$ ) then
12.       $\pi := plan(B, I)$ 
13.    end-if
14.  end-while
15. end while
```

Goal monitoring

In addition to **plan monitoring** above, we can introduce **goal monitoring**.

Basic goal monitoring. Dismiss current plan if:

- It turns out that the intentions (goals) have already been achieved, which we denote *succeeded*(I, B).
- It turns out that the intentions (goals) are impossible to obtain, which we denote *impossible*(I, B).

Version 5 of the agent control loop integrates this type of **goal monitoring**...

The BDI architecture: version 5

Agent Control Loop Version 5

```
1.  $B := B_0; I := I_0;$ 
2. while true do
3.   get next percept  $\rho$ ;
4.    $B := brf(B, \rho); D := options(B, I); I := filter(B, D, I);$ 
5.    $\pi := plan(B, I);$ 
6.   while  $\neg empty(\pi) \wedge \neg succeeded(I, B) \wedge \neg impossible(I, B)$  do
7.     execute( $hd(\pi)$ );
8.      $\pi := tail(\pi);$ 
9.     get next percept  $\rho$ ;
10.     $B := brf(B, \rho);$ 
11.    if not  $sound(\pi, I, B)$  then
12.       $\pi := plan(B, I)$ 
13.    end-if
14.  end-while
15. end while
```

Intention reconsideration

Our BDI-agent only reconsiders its intentions once every time around the outer control loop, i.e., when:

- it has completely executed a plan to achieve its current intentions; or
- it believes it has achieved its current intentions; or
- it believes its current intentions are no longer possible.

In other words, we now have **online (re)planning** but still not **online deliberation**.

Not sufficiently general: Assume my **desire** is to have dinner in London. I might thus form the **intention** of finding a nice and cheap restaurant. I then form a **plan** for searching for such a restaurant. After some hours of unsuccessful search, I might want to reconsider my intention...

Version 6 of the agent control loop integrates **online deliberation**...

The BDI architecture: version 6

Agent Control Loop Version 6

1. $B := B_0; I := I_0;$
2. while true do
3. get next percept ρ ;
4. $B := brf(B, \rho); D := options(B, I); I := filter(B, D, I);$
5. $\pi := plan(B, I);$
6. while $\neg empty(\pi) \wedge \neg succeeded(I, B) \wedge \neg impossible(I, B)$ do
7. $execute(hd(\pi));$
8. $\pi := tail(\pi);$
9. get next percept ρ ;
10. $B := brf(B, \rho); D := options(B, I); I := filter(B, D, I);$
11. if not $sound(\pi, I, B)$ then
12. $\pi := plan(B, I)$
13. end-if
14. end-while
15. end while

The problem of intention reconsideration

Another problem: intention reconsideration is computationally costly!

Dilemma:

- Reconsidering intentions **too seldom**: the agent might continue pursuing its intentions after it should be clear that the intentions are no longer relevant.
- Reconsidering intentions **too often**: the agent may spend too much time reconsidering, and never achieve any of its intentions.

Possible solution: Incorporate an explicit meta-level control component, *reconsider*(I, B), that decides whether or not to reconsider the intentions.

To make any sense, computing *reconsider*(I, B) must be fairly cheap compared to doing the actual reconsideration of intentions.

Version 7 of the agent control loop integrates this...

Agent Control Loop Version 7

```
1.  $B := B_0; I := I_0;$ 
2. while true do
3.   get next percept  $\rho$ ;
4.    $B := brf(B, \rho); D := options(B, I); I := filter(B, D, I);$ 
5.    $\pi := plan(B, I);$ 
6.   while  $\neg empty(\pi) \wedge \neg succeeded(I, B) \wedge \neg impossible(I, B)$  do
7.      $execute(hd(\pi));$ 
8.      $\pi := tail(\pi);$ 
9.     get next percept  $\rho$ ;
10.     $B := brf(B, \rho);$ 
11.    if reconsider( $I, B$ ) then
12.       $D := options(B, I); I := filter(B, D, I)$ 
13.    end-if
14.    if not sound( $\pi, I, B$ ) then
15.       $\pi := plan(B, I)$ 
16.    end-if
17.  end-while
18. end while
```

Summary and final notes on BDI

- In **dynamic (multiagent) environments**, pure (offline) planning is insufficient. The planner must be embedded in a more complex **agent architecture** as e.g. the **BDI architecture**.
- The BDI architecture is centered around the attitudes **belief**, **desire**, and **intention**.
- The BDI architecture can be presented as an **agent control loop**, which can be designed more or less complex depending on the application (version 1–7).
- The BDI architecture is not a fixed algorithm, but a **conceptual model** that can assist the design of agents.
- The origin of the BDI model is a theory of **human practical reasoning** by the philosopher Michael Bratman (1987).

BDI agent application examples

Examples of **implementations** of the BDI architecture:

- The air-traffic management system at one of Sydney's airports.
- At DTU, in a multiagent system consisting of a number of Lego NXT robots solving transportation tasks in a shared environment.
- Some DTU teams in the programming project of this course (using e.g. Jason).